

Optimization: Bitmap Indexes

Knowledge objectives

1. Justify the usefulness of a multi-attribute index
2. Explain what a bitmap index is
3. Justify the usefulness of a bitmap index
4. Explain how to find the record given a bitmap position based on the Hakan factor

Understanding Objectives

1. Decide when a multi-attribute index can be used in the evaluation of a predicate
2. Give the approximate cost of a selection on a table with a bitmap index (without coding)
3. Encode a bitmap index using run-length
4. Evaluate a simple predicate (one logical operation) using run-length encoded bitmap
5. Encode a bitmap index using bitmap-sliced technique
6. Evaluate a simple predicate (one logical operation) using bitmap-sliced technique

MULTI-ATTRIBUTE INDEXES

Usefulness of multi-attribute trees


- ❑ Need more space
 - For each entry, keeps attributes A_1, \dots, A_k
 - May result in more levels, worsening access time
- ❑ Modifications are more frequent
 - Every time one of the attributes in the index is modified
- ❑ It is much more efficient than intersecting RID lists (to evaluate conjunctions)
- ❑ Can be used to solve several kinds of queries
 - Equality of all first i attributes
 - Equality of all first i attributes and range of $i+1$
- ❑ The order of attributes in the index matters
 - We cannot evaluate condition over A_k , if there is no equality for A_1, \dots, A_{k-1}

BITMAP INDEX

B-tree (I)

- ❑ Specially useful for simple queries
 - Without grouping, aggregations, or many joins
- ❑ Works better for very selective attributes (few repetitions per value)
 - Attributes in multidimensional queries are usually not very selective
- ❑ Order of attributes in the index is relevant
 - We can define as many indexes as we want
 - ❑ We can define only one Clustered index
 - ❑ For big tables, they may use too much space

B-tree (II)



Bitmap-index

	Ballpoint	Pencil	Pen	Rubber	A4 paper	A3 paper	Chalk	Eraser
	1	0	0	0	0	0	0	0
	0	0	1	0	0	0	0	0
	0	1	0	0	0	0	0	0
	0	0	0	0	0	0	0	1
	0	0	0	0	1	0	0	0
	1	0	0	0	0	0	0	0
	0	0	0	0	0	1	0	0
	0	0	1	0	0	0	0	0
	0	0	0	0	0	0	1	0
	0	1	0	0	0	0	0	0

	Catalunya	León	Madrid	Andalucía
	1	0	0	0
	1	0	0	0
	0	0	0	1
	0	0	1	0
	0	1	0	0
	1	0	0	0
	0	0	0	1
	0	1	0	0
	1	0	0	0
	1	0	0	0

Querying with bitmaps

SELECT COUNT(*)

...

WHERE articleName IN ['Ballpoint', 'Pencil'] AND region='Catalunya'

Ballpoint

Pencil

Catalunya

1
0
0
0
0
1
0
0
0
0

OR

0
0
1
0
0
0
0
0
0
1

=

1
0
1
0
0
1
0
0
0
1

AND

1
1
0
0
0
1
0
0
1
1

=

1
0
0
0
0
1
0
0
0
1

Updating bitmaps

- Two cases of insertion:
 - Without domain expansion:
 - Add "1"
 - With domain expansion:
 - Add a new vector
- One case of deletion:
 - Change "1" for "0"

Catalunya	León	Madrid	Andalucía	Euskadi
0	0	0	0	0
1	0	0	0	0
0	0	0	1	0
0	0	1	0	0
0	1	0	0	0
1	0	0	0	0
0	0	0	1	0
0	1	0	0	0
1	0	0	0	0
1	0	0	0	0
0	0	1	0	0
0	0	0	0	1

Probabilities with a bitmap

- Probability of a tuple fulfilling P
 SF
- Probability of a tuple NOT fulfilling P
 $1-SF$
- Probability of none of the tuples in a block fulfilling P
 $(1-SF) \cdot (1-SF) \cdot \dots \cdot (1-SF) = (1-SF)^R$
- Probability of some tuple in a block fulfilling P
 $1-(1-SF)^R$

Cost of bitmap per operation

□ Table scan

- $\text{ndist} \cdot \lceil |T|/\text{bits} \rceil \cdot D + B \cdot D$

bits: bits per index block
ndist: different values
v: number of queried values

□ Search for some tuples

- $v \cdot \lceil |T|/\text{bits} \rceil \cdot D + (B \cdot (1 - (1 - \text{SF})^R)) \cdot D$

■ Examples:

- Search for several tuples (given one value): $\text{SF} = 1/\text{ndist}$
 - $\lceil |T|/\text{bits} \rceil \cdot D + (B \cdot (1 - ((\text{ndist} - 1)/\text{ndist})^R)) \cdot D$
- Search for several tuples (given several values): $\text{SF} = v/\text{ndist}$
 - $v \cdot \lceil |T|/\text{bits} \rceil \cdot D + (B \cdot (1 - ((\text{ndist} - v)/\text{ndist})^R)) \cdot D$
- Search for one tuple: $\text{SF} = 1/|T|$
 - $\lceil |T|/\text{bits} \rceil \cdot D + D$

Comparison

- ❑ Better than B-tree for multi-value queries
- ❑ Optimum performance for several conditions over more than one attribute (each with a low selectivity)
- ❑ Orders of magnitude of improvement compared to a table scan (specially for $SF < 1\%$)
- ❑ May be useful even for range queries
- ❑ Sometimes used to manage NULL values
- ❑ Useful for non-unique attributes (specially for $ndist < |T|/100$, i.e. hundreds of repetitions)
- ❑ Bad performance for concurrent INSERT, UPDATE and DELETE
- ❑ Use more space than RID lists for domains of 32 values or more (may be better with compression), assuming uniform distribution and 4 bytes per RID

BITMAP COMPRESSION

Bitmap compression variants

□ Bitmap-encoded index

- Create a look-up table coding values into integers
 - Facilitates coding hierarchies
 - 0 to 3 for Catalan provinces (mask 00xx)
 - 4 to 7 for Galician provinces (mask 01xx)
 - 8 to 15 for Andalucian provinces (mask 1xxx)

□ Bit-sliced index

- Code each value into binary
 - Each bit goes to the corresponding bit-vector
 - We only have $\lceil \log_2 \text{ndist} \rceil$ bit-vectors
- ANDing either the bit or the negated for each column you can check the value:
 - $A \text{ IN } (1,3) \Rightarrow (v_0=1 \text{ AND } v_1=0) \text{ OR } (v_0=1 \text{ AND } v_1=1)$
 $\Rightarrow v_0=1$

Bitmap-encoded-sliced: example

	Barcelona	Tarragona	Lleida	Girona
	1	0	0	0
	1	0	0	0
	0	0	0	1
	0	0	1	0
	0	1	0	0
	1	0	0	0
	0	0	0	1
	0	1	0	0
	1	0	0	0
	1	0	0	0

Encoding



Look-up table

Barcelona	0
Tarragona	1
Lleida	2
Girona	3



Slicing



	2^1	2^0
	0	0
	0	0
	1	1
	1	0
	0	1
	0	0
	1	1
	0	1
	0	0
	0	0

Run-length encoding

- A “run” is a set of i zeros followed by “1”
 - We codify it by a suitable binary encoding of i
- Codification of a run
 - a) Content: Binary coding (e.g., $17 = \overbrace{10001}^{n=5}$)
 - b) Length: $n-1$ ones and a zero (e.g., $5 = 11110$)
- Size gain
 - Size for coding i : $2 \cdot \lceil \log_2 i \rceil$
 - Average value of i : $|T| \cdot (\text{ndist} - 1) / |T| = \text{ndist} - 1$
 - Size of a bitmap: $|T| \cdot 2 \cdot \lceil \log_2 (\text{ndist} - 1) \rceil$



Run-length encoding example

Catalunya	León	Madrid	Andalucía
1	0	0	0
1	0	0	0
0	0	0	1
0	0	1	0
0	1	0	0
1	0	0	0
0	0	0	1
0	1	0	0
1	0	0	0
1	0	0	0

Run-length



Catalunya	León	Madrid	Andalucía
1	3	2	2
0	4	3	2
1	2		2
0	2		3
2			
3			
2			
2			
1			
0			

Binary



Catalunya	León	Madrid	Andalucía
0	1	1	1
0	1	0	0
0	0	1	1
0	1	1	0
1	0		1
0	0		0
1	1		1
1	0		1
1	1		
0	0		
1			
0			
0			
0			

Operating run-length encoded bit-vectors

- We need to decode
 - Not necessarily the whole vector at once
 - One run at a time

A OR B \Rightarrow Being $l_A(i)$ and $l_B(j)$ the length of the runs plus one, we know that the $\Sigma_1^k l_X(i)$ record must be in the output, for X being both A and B

$$\{a | \exists x, a = \Sigma_1^x l_A(i)\} \cup \{a | \exists x, a = \Sigma_1^x l_B(i)\}$$

A AND B \Rightarrow Being $l_A(i)$ and $l_B(j)$ the length of the runs plus one in both indexes, we must generate in the output:

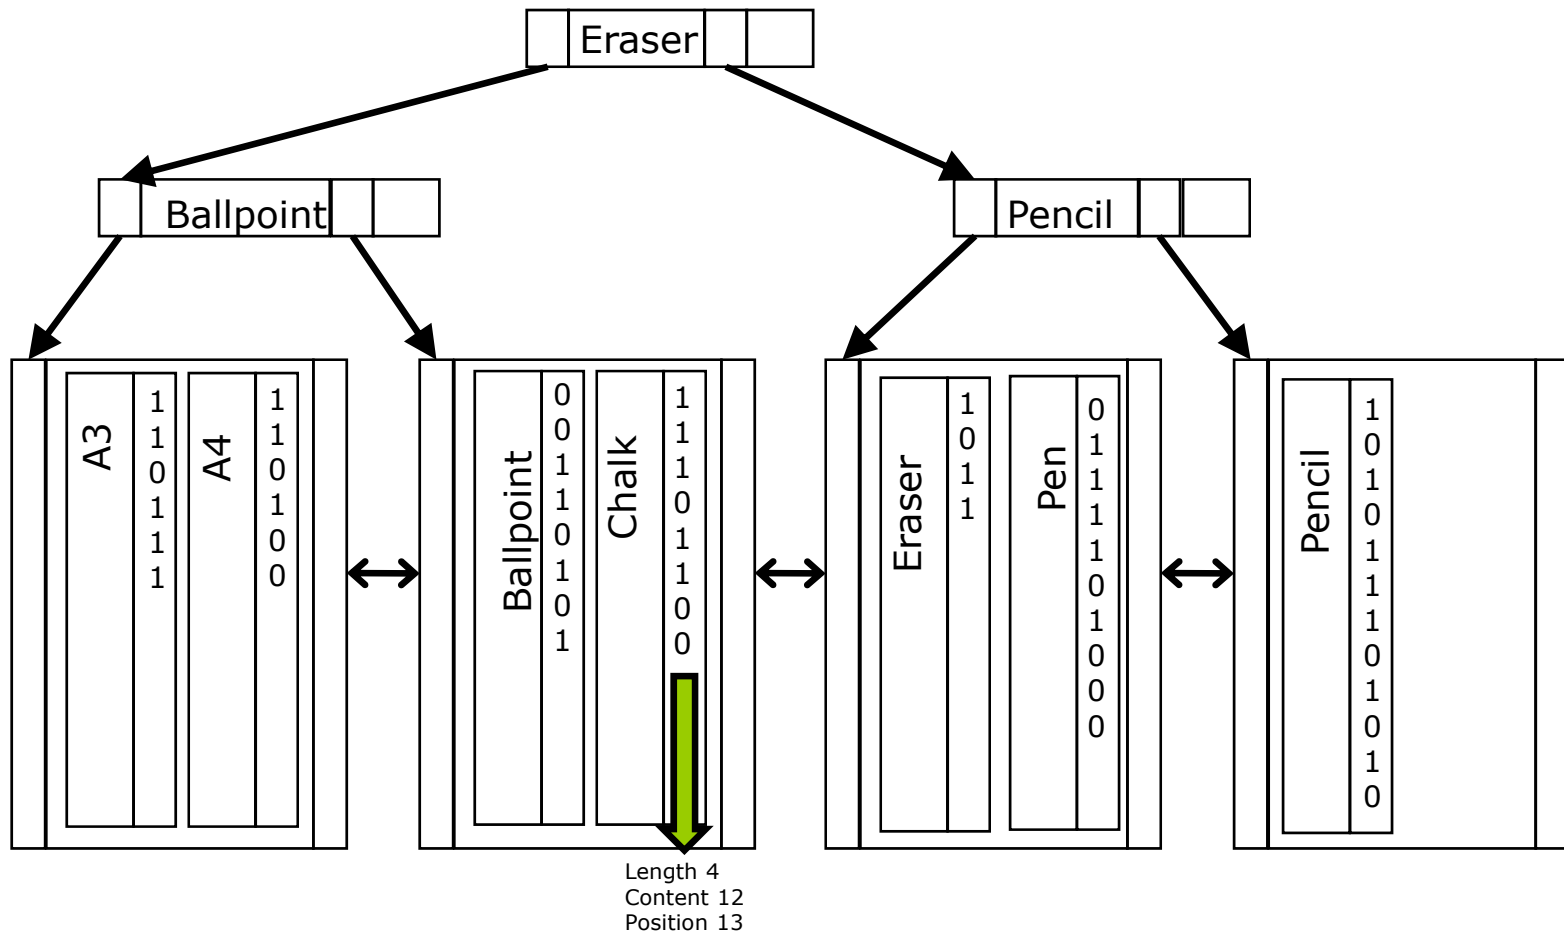
$$\{a | \exists x \exists y, a = \Sigma_1^x l_A(i) = \Sigma_1^y l_B(j)\}$$

Bitmap indexes in Oracle: Declaration

```
CREATE  
[{UNIQUE|BITMAP}] INDEX <name>  
ON <table> (<column>[,column]*);
```


- ❑ Allowed even for unique attributes
- ❑ Does not allow to check uniqueness

Bitmap-index in Oracle: Storage



Finding records of variable length

- a) Maintain a secondary structure
 - Having the position as key
 - Pointing to the disk address of the record
- b) Block-wise implementation
 - Fix the number of bits used per block
 - Use Hakan factor
 - Describes the expected maximum number of rows that a block can hold (e.g., 3)
 - Derived from column definition (and statistics)
 - Affected by
 - Number of columns
 - Column data types and lengths
 - NOT NULL constraints



Setting Hakan factor in Oracle

```
ALTER TABLE <name> [MINIMIZE |  
  NOMINIMIZE] RECORDS_PER_BLOCK;
```

- ❑ The table cannot be empty
- ❑ The table cannot have any bitmap index

CLOSING

Summary

- ❑ Multi-attribute index
- ❑ Bitmap-index
 - Cost
 - Encoding
 - ❑ Bitmap-encoded
 - ❑ Bit-sliced
 - ❑ Run-length encoding

Bibliography

- ❑ P. Valduriez. *Join Indices*. ACM TODS, 12 (2), June 1987. Pages 218-246
- ❑ M. Golfarelli and S. Rizzi. *Data Warehouse Design*. McGraw-Hill, 2009
- ❑ H. Garcia-Molina, J. D. Ullman, and J. Widom. *Database Systems*. Prentice Hall, 2009